

Offer — Research Ops for Regulated Professional Firms

One offer. Two industries (GDPR consultancies + Canadian immigration firms). One landing page variant per industry. Same architecture, same proof, same VSL with two industry-swap intros.

MVO Summary (Vusal's format)

GDPR/privacy-as-a-service firms in Western Europe (10–80 staff) and Canadian immigration boutiques (5–30 lawyers) struggle with consultants and paralegals burning **2–5 hours per matter** manually searching through regulatory corpora (GDPR articles, EDPB guidance, IRCC manuals), past case files, and internal templates — so they would be willing to pay **\$5,000 one-time build + \$800/month** for a RAG-powered research system that lets their team query the full regulatory corpus + the firm's own knowledge base in seconds, with cited, source-linked answers their consultants can hand straight to a client.

Why I strongly believe this works

- **Already built once.** HeyData is the reference architecture. Same FastAPI + Supabase/pgvector + Apify + OpenAI stack. Different corpus, identical pipeline. New client build = 3–4 weeks, not 8.
- **Buyer recognizes themselves immediately.** GDPR firms read the HeyData case study and feel "that's literally my firm." Immigration firms see "regulatory RAG" and translate it without me explaining.
- **Two parallel tests, one piece of work.** I build the offer once, ship two industry landing pages (already have /mortgage as the template), run two Reddit/LinkedIn outreach motions in parallel. Whichever industry hits first becomes the niche.
- **Price is below enterprise radar but real revenue for me.** \$5K + \$800 means I close 3 builds and I'm at \$15K cash + \$2.4K MRR. That's the runway to invest in the second proof point.

Checklist

- Solves a painful problem — YES (validated by HeyData paying for it)
 - Is buildable — YES (already shipped once in production)
 - Won't take too long — YES (3–4 weeks per client, 6 weeks max for the harder one)
-

Grand Slam Offer Breakdown

Step 1 — Starving Crowd Score

Criterion	GDPR consultancies	Canadian immigration
Massive pain	9/10 — same regulatory questions answered hundreds of times across clients	8/10 — IRCC manuals are 800+ pages, paralegals waste days
Purchasing power	9/10 — €150-250/hr billable rates, \$4-5K is a half-day of consultant time	7/10 — boutique immigration firms run leaner than corporate firms but have real revenue
Easy to target	10/10 — IAPP directory, LinkedIn "DPO as a service" filter, the field is small	8/10 — CBA Immigration Section + CICC + LinkedIn Sales Nav
Growing	9/10 — GDPR enforcement up, AI Act layered on top, more compliance burden every quarter	9/10 — Canadian immigration volume hit record highs, IRCC backlogs are political news

Verdict: Both crowds qualify. GDPR is slightly stronger because the buyer is already pre-validated by HeyData. Immigration is the better long-term niche because there's near-zero AI vendor competition.

Step 2 — Value Equation Scorecard

Lever	Score	How the offer fixes it
Dream outcome	Strong	"Get back 10+ hours/week per consultant. Answer client questions with citation-backed precision in seconds, not hours."
Perceived likelihood	Medium → fixed by HeyData full case study + live demo on their own docs	Fixed via Diagnostic Demo bonus (see stack)

Lever	Score	How the offer fixes it
Time delay	Weak (8 weeks felt long) → fixed by 4-week sprint guarantee + first usable build at week 2	Compressed via "Week 2 Working Build" milestone
Effort & sacrifice	Strong if delivered done-for-you, weak if buyer has to wrangle docs	Fixed by Document Onboarding Concierge (we ingest, chunk, validate — they hand over a folder)

Step 3 — Problems → Solutions

Buyer problem	Solution baked into offer
"I don't know if RAG is accurate enough for regulatory work"	Live demo on their own docs in 48 hours before they pay anything
"Will this hallucinate and embarrass my firm?"	Citation-required output mode — every answer links to source paragraph
"We've been burned by consultants who disappear after delivery"	Fixed-price 4-week sprint + 30-day post-launch hypercare
"What if my team won't adopt it?"	90-min team training session + Slack support channel for first 60 days
"Our docs are messy and half-scanned"	Document Onboarding Concierge — we OCR, clean, structure
"What about IRCC/GDPR updates?"	Monthly regulatory refresh built into retainer
"Is my data safe?"	Self-hosted option (their Supabase, their OpenAI key) — they own everything
"I don't want vendor lock-in"	Source code delivered at end of engagement
"I can't justify \$5K to my partners"	Half-day audit deliverable that quantifies their hours wasted before the build
"What if I want to add a new doc type later?"	First 5 doc types ingested free; ongoing ingestion via retainer

Buyer problem	Solution baked into offer
"We already use ChatGPT"	Side-by-side accuracy benchmark on their docs (will show 60-80% better citation accuracy)
"Our partners are skeptical of AI"	Demo built specifically for the most-skeptical partner in the firm

Step 4 — The Stack

Core deliverable: Custom RAG research system, deployed in production, on their stack or yours.

Stack item	What they get	Anchor value
The Build — production RAG system on their corpus	FastAPI backend, Supabase/pgvector, OpenAI embeddings, citation-linked retrieval. Yours forever.	\$15,000
Diagnostic Demo (pre-build)	48-hour proof: I run the pipeline on a sample of their docs and show them retrieval working before they sign	\$1,500
Document Onboarding Concierge	I ingest, OCR, chunk, and validate their entire corpus. They hand over a folder, I hand back a working system.	\$2,000
Custom citation UI — branded, embeddable	Searchable chat interface with one-click source verification, branded to their firm	\$3,000
Team Training Session — 90 minutes, recorded	Live walkthrough for their consultants/paralegals + recording for new hires	\$800
30-Day Hypercare	Direct Slack/email access, daily standups for week 1, weekly thereafter, all bug fixes free	\$2,400

Stack item	What they get	Anchor value
Source code delivery	Full repo, deployment docs, no vendor lock-in	\$2,000
First 5 document types ingested free	Adapters for any 5 file types (PDF, DOCX, scanned PDFs, web-scraped, structured exports)	\$1,000
Stack total		\$27,700
Your price		\$5,000 build + \$800/month

The \$800/month covers: monthly regulatory refresh (new EDPB opinions / IRCC updates), uptime monitoring, model updates, and 4 hours of system tuning.

Step 5 — Enhancers

Scarcity: "I take 2 new clients per month. I'm currently building for [X]. Next available slot: [month]." Real, because solo founder.

Urgency: Founding-cohort pricing. The first 5 firms get \$5K + \$800. After that, it's \$10K + \$1,500 (the original MVO pricing). This is honest urgency — the price genuinely does go up once case studies exist, because the risk to the buyer drops.

Bonuses (stacked into the offer above): Diagnostic Demo, Document Onboarding Concierge, Team Training, 30-Day Hypercare, Source Code Delivery. Each is a real friction-killer, not filler.

Guarantee — the Working Build Guarantee:

"If you don't have a working research system answering questions on your own documents with cited sources by day 28, I refund every dollar of the build fee and you keep the source code I've delivered up to that point."

This is conditional, performance-based, and specific. It works because: (a) I've built it before, so the timeline risk is real but manageable, (b) the source code clause means they're never empty-handed, (c) it forces me to ship.

Naming — MAGIC formula applied:

- Magnetic reason why: get back hours, kill the manual ctrl+F search
- Avatar: regulated professional firms drowning in their own document corpus

- Goal: instant cited answers across the firm's knowledge
- Interval: 30-day deployment
- Container: a productized build sprint

3 candidate names:

1. **Research Ops** — clean, category-defining, positions the system as infrastructure not a feature. Easy to extend across verticals.
2. **The 30-Day Research System** — pure outcome + timeline name. Very Hormozi.
3. **CitationDesk** — branded SaaS-feeling name, harder to extend across verticals.

Pick: Research Ops. It works for both GDPR firms ("Research Ops for compliance consultancies") and immigration firms ("Research Ops for immigration boutiques"). It's the same product, repositioned per vertical. Extensible to estates, mortgage, PI later without renaming.

Step 6 — Final Pricing

Founding cohort (first 5 clients): \$5,000 build + \$800/month Standard pricing (clients 6+): \$10,000 build + \$1,500/month

Pricing rationale:

- \$5K is below the threshold where partners need committee approval at boutique firms
- \$5K is one consultant's billable week — a no-brainer if it returns 10+ hours/week to the team
- \$800/month is below the "is this a real line item" threshold for firms that already pay \$200-500/month for legal research tools
- The price doubles after 5 case studies because the risk profile changes — buyer 6 isn't taking a chance on a vendor with one reference, they're buying a proven system
- Margin: at \$5K + 4 weeks of work, that's ~\$1,250/week of focused build time.
Sustainable for now, not forever. Honest signal to the buyer that this is founding pricing, not a permanent rate.

Marketing Plan

Channels (ranked by expected ROI)

1. **LinkedIn outreach** — handwritten, HeyData-anchored, ICP-precise. 30 messages/day across both verticals.

2. **Reddit content marketing** — same playbook that worked for the early NovraAI leads. r/PrivacyTools, r/gdpr, r/ImmigrationCanada, r/LawCanada. Builder voice: case studies, technical breakdowns, no pitching.
3. **Industry directories** — IAPP (GDPR), CBA Immigration Section + CICC (immigration). Cold email where allowed.
4. **Twitter/X** — once the suspension resolves. Build-in-public posts about the second client. Slow burn.

Outreach hooks — vertical-specific

GDPR / Privacy consultancies:

"I built a RAG research system for HeyData (German GDPR compliance firm) — they query the full GDPR corpus + EDPB guidance + their own templates and get cited answers in seconds. Saving their consultants 10+ hours/week. Same architecture works for any privacy-as-a-service firm. Worth a 15-minute walkthrough?"

Canadian immigration:

"Your paralegals spend half a day digging through IRCC manuals to assemble background for an LMIA application. I built the same kind of system for a German GDPR consultancy — they ask their full regulatory corpus a question, get a cited answer in 10 seconds. Identical architecture works for IRCC + your case files. Want to see it running on your docs?"

Proof mechanism

The Diagnostic Demo is the conversion lever. It's not a slide deck or a video — it's me running the pipeline on 50-100 of their actual documents over 48 hours and sending back a video showing real questions answered with real citations. Zero risk to them, high-signal commitment from me. Whoever requests one is 70%+ likely to close.

First-week action checklist

- ☐ Spin up two industry landing pages: /gdpr (clone of /mortgage with GDPR copy) and /immigration (same)
- ☐ Replace the case study card on each page with the full HeyData case study (name, metrics, what was built)
- ☐ Build a 60-second Loom of HeyData's system answering 3 questions live — this is the proof asset for outreach
- ☐ Pull 50 GDPR consultancy contacts from IAPP + LinkedIn ("DPO as a service" search, 10-80 employees, EU)

- ☐ Pull 50 Canadian immigration contacts from CBA + LinkedIn (Managing Partner, 5-30 lawyers, Toronto/Vancouver/Montreal)
- ☐ Send 20 messages/day to each list, 5 days, track open + reply
- ☐ Whichever vertical replies more becomes the focus for the next 30 days

What I'd cut from the original 4 MVOs

- **PI demand letter generator** — different product (generation, not retrieval), heavier competition (EvenUp, Eve, Supio are all funded), HeyData proof doesn't transfer. Park it. Revisit after 5 RAG clients are paying.
- **Estates/probate** — genuine white space but slow buyer and harder build (state-by-state statute variance is real engineering, not a copy-paste). Park as the SaaS play once you have the agency cash flow to fund it.

Risk to flag honestly

- One proof point is thin. Two of the first three sales calls will ask "do you have a Canadian immigration reference?" or "do you have another GDPR firm?" The answer for now is "you'd be the second / third — that's why pricing is half the standard rate." That answer works for buyer 1-3, then stops working.
- The \$800/month retainer is too cheap for what's promised. Once 3 clients are live, the standard pricing of \$1,500/month should become the floor. The \$800 is acquisition pricing, not unit economics.
- Solo capacity ceiling at 2 builds/month is real. By client 5, you need a contractor or you're capped at \$10K/month MRR.