VSL — Research Ops

Length: ~10 minutes **Offer Type:** AI Automation / Productized Service ($5K build + $800/mo retainer) **Target Viewer (ICP):** Founder/Managing Partner of a GDPR consultancy (10-80 staff) OR Canadian immigration boutique (5-30 lawyers) **Price + Anchor:** $5,000 build + $800/month, anchored against $42,000-$80,000 of consultant/paralegal time wasted on manual research per year **CTA:** Book a 15-minute strategy call → novraai.dev

> **How to use this script:** Beats 1, 2, and the case study in Beat 5 are the only sections that change between the GDPR and Immigration versions. Everything else is identical. Film the master version once, reshoot 30 seconds at the top + one story card. That's it.

Beat 1: The Hook `[0:00 – 0:30]`

`[VISUAL: Vusal on camera, tight crop, plain background. Bold on-screen text: "Your team is spending €X/year doing ctrl+F." Number animates in.]`

(GDPR version): Your firm has a problem you've stopped seeing because it's been there since day one. A senior consultant on €180 an hour is sitting in front of a 600-page PDF of EDPB guidance, scrolling. Looking for one paragraph. They'll find it in fourteen minutes. They'll do that six more times today. Across your team, that's somewhere between forty and eighty thousand euros a year of billable time spent on ctrl+F. Stay with me — I'm going to show you exactly how to get that back.

(Immigration version): Your paralegal is sitting in front of an 800-page IRCC operational manual right now. They're scrolling. They're looking for one paragraph that satisfies a condition on an LMIA file. They'll find it in fourteen minutes. They'll do that six more times today. Across your team, that's somewhere between forty and eighty thousand dollars a year of billable hours spent on ctrl+F. Stay with me — I'm going to show you exactly how to get that back.

Beat 2: The Problem `[0:30 – 1:45]`

`[VISUAL: B-roll – overhead shot of someone scrolling a long PDF, a calculator overlay running in the corner adding hours, a stack of` binders, a Slack thread of someone asking "anyone remember which article covers this?"]

Let's do the math out loud, because I think you'll find this number is bigger than you think.

You have, let's say, four consultants. Each one spends — conservatively — two hours a day searching documents. Searching the regulatory corpus. Searching old client files for precedent. Searching internal templates for the right paragraph. Two hours, every day, every consultant.

That's eight hours a day across the team. Forty hours a week. **Two thousand hours a year.** At your blended billable rate, that's somewhere between forty and a hundred thousand a year of capacity that doesn't generate a single dollar of client revenue.

[VISUAL: Big on-screen number — "2,000 hours/year" — then animates into a dollar figure.]

And here's the part that hurts more. Half of those searches are for questions you've already answered. Last year. Last month. Last Tuesday. The answer exists, somewhere, in your own files. Your team just can't find it fast enough so they redo the work.

[VISUAL: Side-by-side — a screen full of folders labeled by client/year, vs. a single search bar.]

You've thought about hiring another paralegal. The math doesn't work — they'd be searching documents too. You've tried ChatGPT. It hallucinates citations and your partners will not let that near a client deliverable. You've thought about a knowledge management system. The last one a vendor sold you took six months to deploy and nobody used it.

So your team keeps scrolling.

Beat 3: The Reframe [1:45 — 2:45]

[VISUAL: Tone shift — brighter footage, Vusal back on camera. "What if..." appears as on-screen text.]

What if your team could ask a question — in plain English, the way they'd ask the senior partner — and get an answer in eight seconds, with the exact paragraph from the exact document highlighted, with a one-click link to the source?

Not a summary. Not a chatbot guess. The actual paragraph. From the actual document. Every time.

[VISUAL: Screen recording — search bar, question typed, citation pops up, click expands the source PDF to the highlighted paragraph.]

Your consultant types "what's the latest EDPB position on legitimate interest for B2B marketing" — eight seconds later they have the article, the EDPB opinion, and the paragraph from your firm's own template that handles it. They paste it into the client email. They move to the next file.

That fourteen-minute search becomes a thirty-second confirmation. The senior consultant stops being your team's search engine. The associate stops being afraid to ask. The paralegal handles cases that used to require partner review.

The work doesn't get faster. The work that doesn't need to happen — disappears.

Beat 4: The Credibility [2:45 – 4:00]

[VISUAL: Vusal on camera, lower-third with name + title: "Vusal Novruzov, Founder, NovraAI." Logo wall is sparse — HeyData logo prominent, small "Built in Tbilisi" tag.]

I'm Vusal Novruzov. I'm an AI engineer and the founder of NovraAI. I build research systems for regulated professional firms — companies where every answer has to be cited, every document has to be traceable, and "the AI made it up" is a fireable offense.

I'm not a generalist who pivoted into AI when ChatGPT came out. I've been in production with retrieval-augmented systems for years. My stack is FastAPI, Supabase with pgvector, Apify, OpenAI embeddings — the same stack that powers research tools at companies you'd recognize. The difference is I build it for one firm at a time, custom to their corpus, owned by them.

[VISUAL: Cut to HeyData case study card — logo, location, what they do.]

The system I'm about to show you was built first for a company called HeyData. They're a German GDPR compliance-as-a-service firm. They manage hundreds of client accounts across the EU. Their consultants were drowning in the same problem your team is in right now — same regulatory corpus, same internal templates, same repeated questions. We deployed the system. Their consultants now query the full GDPR corpus, EDPB guidance, and their own internal templates from one search bar. Research time dropped from hours to minutes.

```
[VISUAL: Screen recording of HeyData's system answering a real GDPR question, with citation expanding to source.]
```

That's the proof. One firm. One deployed system. Real consultants using it daily. Everything I'm going to tell you in the next five minutes is built on that.

Beat 5: The Proof `[4:00 – 5:15]`

```
[VISUAL: Story cards for each example — name, screenshot of the system, before/after numbers.]
```

Story 1: HeyData — the reference build

```
[VISUAL: HeyData logo, what they were before, what changed.]
```

Before we built the system, HeyData consultants were searching across the full GDPR text, EDPB guidance documents, national DPA rulings, and their own internal compliance templates manually. A single client query — "does this processing activity require a DPIA under Article 35?" — could take a senior consultant 20 to 40 minutes to research properly with citations.

After deployment: same question, same precision, eight seconds. The consultant validates the citation, sends the client answer. Their team handles a higher caseload with the same headcount. The system runs on their stack — they own it. We maintain it.

Story 2: The vertical adapts cleanly

```
[VISUAL (GDPR version): Demo of the system loaded with a different EU privacy firm's templates — anonymized.] [VISUAL (Immigration version): Demo of the system loaded with IRCC operational manuals + sample case file — show a "what are the requirements for an LMIA exemption under C-16" query.]
```

(GDPR version): The architecture isn't tied to GDPR. We've loaded the system with EU AI Act guidance, with DORA documentation, with ISO 27001 templates. Every regulatory corpus is just another set of documents. The pipeline doesn't care.

(Immigration version): The architecture isn't tied to GDPR. The same pipeline ingests IRCC operational manuals, program delivery instructions, and your firm's past case files. We've tested it on a sample IRCC corpus — same retrieval quality, same citation precision. The pipeline doesn't care which regulator wrote the document.

Story 3: The build risk is gone

[VISUAL: Calendar timeline graphic — Week 1, 2, 3, 4 with milestones.]

The first build took us months because we were figuring it out in the dark. Every build after HeyData runs on the same architecture, the same ingestion pipeline, the same citation framework. New corpus, new firm, three to four weeks. That's not a sales pitch — that's because the engineering risk has already been paid for.

Beat 6: The Offer Reveal [5:15 – 6:00]

[VISUAL: Big logo/name reveal — "Research Ops" — animates in. Tagline below: "Cited answers. Your documents. 30 days."]

The offer is called **Research Ops**.

It's a custom RAG research system, deployed in production on your firm's document corpus, in 30 days, with every answer citation-backed and source-linked.

It's not a SaaS subscription. It's not a chatbot you bolt onto your website. It's a research operations system built for your firm specifically — your regulations, your templates, your client files, your team's vocabulary. You own the source code. You own the deployment. We maintain it.

One offer. One sprint. One system that gives your consultants their hours back.

Beat 7: The Mechanism [6:00 – 8:30]

[VISUAL: Component cards animating in one at a time. Each component has a small demo or diagram.]

Here's exactly what gets built and shipped to you in 30 days. Six components.

1. The Diagnostic Demo (week 0, before you pay anything)

[VISUAL: A short video frame — Vusal on a Loom with a few of the buyer's documents loaded, asking 3 questions, citations popping up.]

Before you commit, I take 50 to 100 of your actual documents and run the pipeline on them. Forty-eight hours later you get a recorded video showing real questions answered with real citations from your real documents. If retrieval quality isn't where you need it — we don't move forward and you've spent zero dollars.

2. The Document Onboarding Concierge (weeks 1-2)

`[VISUAL: Diagram — folder full of mixed document types flowing into a clean, structured corpus.]`

You hand over a folder. PDFs, Word docs, scanned files, email exports, structured data — whatever shape your knowledge is currently in. I OCR the scanned files, chunk every document semantically, validate retrieval quality on a sample, and load it into your vector database. Five document types covered in the base build. Your team doesn't lift a finger.

3. The Retrieval Engine (weeks 2-3)

`[VISUAL: Architecture diagram — query → vector search → reranker → citation builder → answer with source link.]`

Custom RAG pipeline tuned to your corpus. Semantic chunking, hybrid retrieval, reranking, and citation extraction. Built on FastAPI and Supabase pgvector — boring, proven, fast. Every answer comes with a source paragraph and a one-click link to the underlying document. No hallucination — the system literally cannot return an answer without a source.

4. The Branded Search Interface (week 3)

`[VISUAL: Screenshot of a clean chat interface, branded with a fictional firm's logo, citation panel on the right.]`

A clean, branded chat interface your consultants actually use. Search bar, conversation history, citation panel, source preview. Embeddable into your existing tools or hosted standalone at research.yourfirm.com. Designed to look like your firm built it, not like a generic AI tool.

5. The Team Training (week 4)

`[VISUAL: Zoom call thumbnail with a team of 6 consultants, recorded session.]`

90-minute live training session for your team. How to phrase queries, how to validate citations, how to spot edge cases. Recorded for new hires. Plus a Slack channel with direct access to me for the first 60 days — every question, every bug, every "can it do this" — answered same day.

6. The Monthly Refresh (ongoing, $800/month)

[VISUAL: Calendar icon ticking through months, regulations being added.]

Your regulatory landscape changes. New EDPB opinions. New IRCC bulletins. New internal templates. Once a month, I refresh the corpus, retune retrieval, and patch any quality drift. Plus uptime monitoring and four hours of system tuning. Your system stays accurate without your team thinking about it.

That's Research Ops. Six components. Thirty days. One sprint, one system, your firm owns it.

Beat 8: The Math [8:30 – 9:15]

[VISUAL: Comparison table on screen. Three columns: "Hire a paralegal," "Enterprise legal AI," "Research Ops." Numbers stack dramatically.]

Let's price the alternatives.

Option A — hire another paralegal to do more of the searching. Fully loaded cost: $65,000 to $90,000 a year, plus they're capped at 40 hours a week, plus they need vacation, plus they search at human speed.

Option B — buy enterprise legal AI like Harvey or CoCounsel. Floor price for a firm your size: $80,000 to $200,000 a year. Generic across all law. Doesn't know your templates. Doesn't know your case history. Locks your data into their system.

Option C — Research Ops. Five thousand dollars one-time. Eight hundred dollars a month. Built specifically on your corpus. Source code yours. Deployed in 30 days.

[VISUAL: Big number reveal — $5,000 build, $800/mo. Then the delta vs. Option A and B animates in.]

That's a tenth of the price of the alternatives. Built specifically for your firm. Owned by you. The math isn't close.

And remember the number from the beginning — your team is currently burning forty to a hundred thousand a year on manual document search. The system pays for itself the first month it's deployed. Every month after that is pure recovered capacity.

Beat 9: The Objections `[9:15 – 10:00]`

`[VISUAL: Vusal on camera, direct. Each objection appears as on-screen text and gets crossed out as it's addressed.]`

Objection 1: "Will this work for my specific corpus?"

Probably yes — but you don't have to take my word for it. The Diagnostic Demo runs on your actual documents before you spend a dollar. If retrieval quality isn't where it needs to be on your corpus, we both find out in 48 hours and we don't move forward. Zero risk to you.

Objection 2: "What if it hallucinates and my consultants embarrass the firm?"

The system is built so it cannot return an answer without a source paragraph and a citation link. It's not a generative chatbot pretending to be a researcher. It's a retrieval system that surfaces the exact text from your documents. Your consultants validate the citation in one click. They're never reading machine-generated regulatory advice — they're reading your documents, found faster.

Objection 3: "I'm not technical. Is my team going to be able to use this and maintain it?"

If your team can use Google, they can use this. The interface is a search bar. That's it. And the maintenance — corpus refreshes, model updates, monitoring — is what the $800 monthly retainer covers. You don't touch it.

Objection 4: "Is this just another vendor lock-in trap?"

No. You get the source code. You get the deployment documentation. The system runs on your infrastructure or mine — your choice. If you ever want to take it in-house, the handover is built into the engagement. I'm selling you a system you own, not a subscription you can't leave.

Guarantee `[10:00 – 10:20]`

`[VISUAL: Big guarantee badge on screen. Vusal direct to camera.]`

And here's how we both stay aligned. **The Working Build Guarantee.** If by day 28 you don't have a working research system answering questions on your own documents with cited sources, I refund every dollar of the build fee. And you keep the source code I've delivered up to that point. I take the risk. You don't.

Beat 10: The Close `[10:20 – 10:50]`

`[VISUAL: Vusal on camera, calm, direct. CTA URL on screen for the entire close — novraai.dev. End frame: CTA + "2 build slots / month."]`

Status check

Every week your firm doesn't have this system, your team is spending another 40 hours scrolling through documents your AI system could surface in seconds. Forty hours that don't bill. Forty hours your competitors who deployed this system are spending on actual client work.

Clear instruction

If you want to see this running on your firm's documents, book a 15-minute strategy call at novraai.dev. We'll talk through your corpus, the questions your team gets asked most, and whether Research Ops is the right fit. If it is, the next step is the Diagnostic Demo on your documents — at zero cost.

Urgency

I take two new clients per month. Founding-cohort pricing — five thousand build, eight hundred a month — applies to the first five firms only. After that, it's ten thousand build and fifteen hundred a month, because the case studies will exist and the risk profile changes.

`[VISUAL: Final frame — "novraai.dev" + "2 slots / month" + "Founding pricing for the first 5 clients"]`

Book the call. Get your hours back.

Alternate Hook Openings

Three swap-in openings for A/B testing. Each replaces Beat 1 only (`[0:00 – 0:30]`).

Alt Hook 1 — The bold claim with a number

[VISUAL: Vusal direct to camera. Bold on-screen text: "Your team loses 2,000 hours/year to ctrl+F."]

Your firm is losing two thousand billable hours a year to ctrl+F. Not to bad cases. Not to slow clients. To your team scrolling through PDFs, looking for one paragraph, one article, one precedent. I'm going to show you the exact system one of my clients uses to get those hours back — built in 30 days, owned by them, citation-backed every time. Watch the next eight minutes. If the math doesn't make sense, close the tab.

Alt Hook 2 — The diagnostic question

[VISUAL: Vusal leaning in to camera. On-screen text: "When was the last time a consultant on your team..."]

When was the last time a consultant on your team spent 20 minutes searching for a paragraph they knew existed somewhere in your firm's files? Yesterday? This morning? Right now? If your team is searching documents more than they're answering clients, the bottleneck in your firm isn't capacity. It's retrieval. And there's a system that fixes it in 30 days. I built it. I'm going to show you exactly how it works.

Alt Hook 3 — The contrarian truth

[VISUAL: Vusal direct, slight smirk. On-screen text: "You don't need another paralegal."]

You don't need another paralegal. You don't need Harvey. You don't need a 200,000-dollar enterprise AI subscription. What your firm actually needs costs five thousand dollars, takes 30 days to deploy, and turns your existing team into the most efficient research operation in your vertical. Most consultancies your size are about to spend the next two years getting sold expensive AI tools that don't fit. Watch this video before you write that check.